

---

# Diversify Your Vision Datasets with Automatic Diffusion-Based Augmentation

---

**Lisa Dunlap Alyssa Umino Han Zhang Jiezhi Yang Joseph E. Gonzalez Trevor Darrell**  
UC Berkeley

## Abstract

Many fine-grained classification tasks, like rare animal identification, have limited training data and consequently classifiers trained on these datasets often fail to generalize to variations in the domain like changes in weather or location. As such, we explore how natural language descriptions of the domains seen in training data can be used with large vision models trained on diverse pretraining datasets to generate useful variations of the training data. We introduce ALIA (Automated Language-guided Image Augmentation), a method which utilizes large vision and language models to automatically generate natural language descriptions of a dataset’s domains and augment the training data via language-guided image editing. To maintain data integrity, a model trained on the original dataset filters out minimal image edits and those which corrupt class-relevant information. The resulting dataset is visually consistent with the original training data and offers significantly enhanced diversity. On fine-grained and cluttered datasets for classification and detection, ALIA surpasses traditional data augmentation and text-to-image generated data by up to 15%, often even outperforming equivalent additions of real data. Code is available at <https://github.com/lisadunlap/ALIA>.

## 1 Introduction

While modern pretraining data are incredibly diverse, datasets for specialized tasks such as fine-grained animal identification are often much less so, resulting in trained classifiers that fail when encountering new domains such as a change in weather or location. An effective method to address this is to add more training data from the test domain [37], but obtaining and labeling this additional data is often costly and it can be challenging to determine which domains to gather data from.

To address this, recent works have utilized image generation models trained on large pretraining datasets to supplement incomplete training sets by generating diverse examples from generative models fine-tuned on the training set [8, 14, 1]. Furthermore, previous work in language-guided data augmentation with vision and language models relies on user-supplied domain descriptions [7] or descriptions generated from word-to-sentence models [9].

Although these methods can be effective in increasing image diversity, they either require finetuning the image generation model, which can be prohibitively expensive, or generating images which have no visual grounding in the original training data. While the latter may be suitable for common benchmarks, such as ImageNet [6], it proves to be much less effective when we move to a specialized setting where generative models like Stable Diffusion [31] cannot recreate images which resemble the training data from text alone, as shown in Figure 1. In our work, we focus on how to utilize pretrained vision and language models for image captioning and generation as a *translation* layer between task-specific image data and task-agnostic natural language descriptions of domains. Since these high-level domain descriptions are well-represented by image generation models like Stable Diffusion, we can use them to perform *language-guided image editing* of the specialized training data.

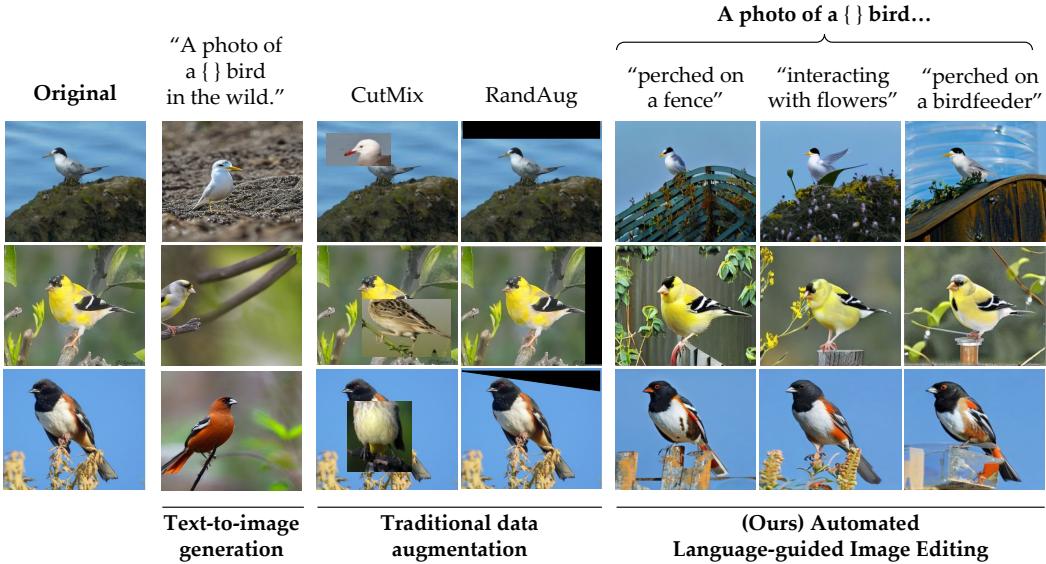


Figure 1: **Overview.** Example augmentations using text-to-images generation, traditional data augmentation methods, and our method, Automated Language-guided Image Editing (ALIA) on CUB [41]. Images generated by ALIA retain task-relevant information while providing more domain diversity as specified by the prompts. Within the prompts, { } indicates the specific class name.

This produces images which are visually consistent with the training data, vary the task-agnostic domains, and preserve the task-relevant information present in the original image.

Specifically, our method ALIA (Automated Language-guided Image Augmentation) first generates captions for each image, summarizes the captions into a short list of domain descriptions with a large language model (LLM), and then uses these descriptions to generate edits of the training data with Stable Diffusion. To ensure data quality, we use a classifier trained on our original training set to remove images which (1) do not change the domain of the original image as desired or (2) corrupt task-relevant information. After filtration, we are left with an edited dataset visually consistent with the original data and representing all domains seen in testing (Figure 1). ALIA does not require finetuning the image captioning or image generation model, nor does it require user-supplied prompts.

We evaluate on fine-grained bird classification (CUB [41]), domain generalization (iWildCam [17]), and contextual bias (Planes [22]) datasets. We show that the addition of ALIA generated data outperforms traditional data augmentation techniques and text-to-image generated data by up to 15%, even beating the performance of adding in real data on iWildCam. Furthermore, we investigate how our domain descriptions produce more useful edits than user-provided prompts, and examine the effect of filtering, and the choice of image editing techniques (Section 5.4).

## 2 Related Works

**Supplementing Training Data with Generative Models.** Using generative models for data augmentation is a well-explored area of research, specifically in medicine [8, 33], domain adaptation [12] and bias mitigation [34]. These methods train or use a pretrained GAN to generate images from the desired distribution. Furthermore, GANs have been used to supervise dense visual alignment [29] and generate pixel-level annotations from a small amount of labels [45, 18, 39]. Recently, several works [9, 35, 38] have shown diffusion models’ ability to generate training data in zero or few shot settings as well as generate hard training examples [13]. While these works do show the promise of diffusion-generated data, models trained on diffusion-generated data obtain significantly worse accuracy than models trained on real datasets unless finetuned for that specific task [1, 38]. In contrast, we use diffusion models to do *image editing* with text rather than generating images from text alone, resulting in augmentations that closely resemble the training data without finetuning.

**Traditional Data Augmentation.** Traditionally, data augmentation techniques involve random flipping, cropping, color shifting, etc. to manipulate the original images and create new versions of these images [36]. More recently proposed mixup-based data augmentations aim to make the augmented data more diverse by cutting and pasting patches of two input images [42], applying convex combinations [43], and adaptively learning a sample mixing policy [10, 20, 15, 5]. These methods often result in images that look unnatural as shown in Figure 1, while our method aims to create augmentations which are visually consistent with the original training data.

**Language-guided Data Augmentation for Classification.** Recent works [7, 9, 30] have explored how natural language descriptions of the domains and classes seen during testing can be used to improve performance and generalization in a few-shot or domain adaptation setting. These methods typically employ text-to-image generation with prompts sourced from language models [9], or augment training images using user-provided descriptions of both training and unseen test domains in a shared vision and language space [7]. Similar to our work, these techniques use the high-level domain knowledge of large models to diversify image data with language, but crucially they either rely on user-provided domain descriptions or generate descriptions and augmentations that are not grounded in any real image data. In contrast, the foundation of our work is to ground both the domain descriptions and the generated training data in the given training set.

### 3 ALIA: Automated Language-guided Image Editing

Given a labeled training dataset, we aim to augment the dataset with images that are edited to improve the representation of various *domains*. We define *domain* to be any aspect of an image that is not intended to be used for classification (e.g. location, weather, time of day). The key insight of our method is to utilize image captioning and image generation models trained on large amounts of pretraining data to summarize the domains in the training set and use those descriptions to augment training data using text-conditioned image editing.

Our method consists of 3 stages: generating domain descriptions, generating useful image edits, and filtering out poor edits. An overview of our method is shown in Figure 2.

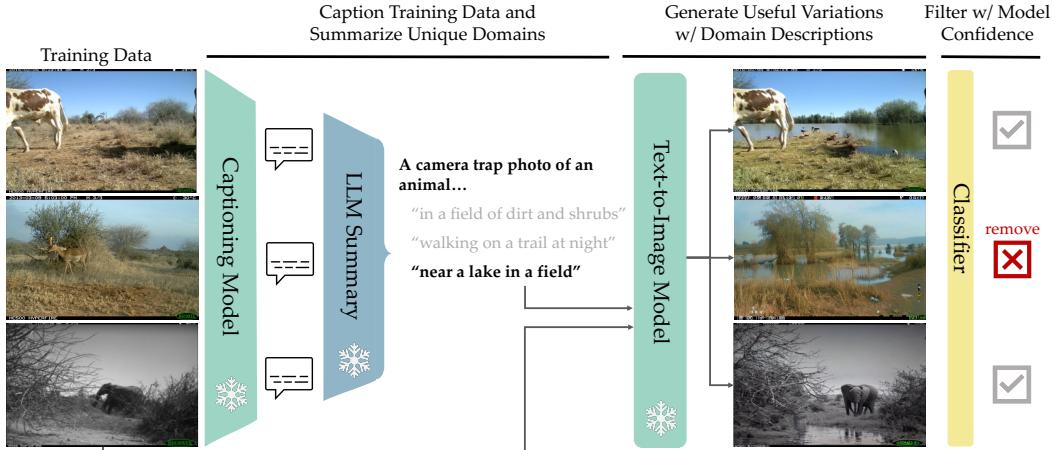


Figure 2: **ALIA.** Given a specialized dataset, we caption all images in our dataset using a pretrained captioning model, and feed these captions into a large language model to summarize them into a small (<10) set of natural language descriptions. Utilizing these descriptions, we perform text-guided image augmentation via a pretrained text-to-image model, thereby generating training images that align with the described settings. Finally, we apply two filtering techniques: a CLIP-based semantic filter to eliminate obvious edit failures, and a confidence-based filter that removes more subtle failures by filtering edits confidently predicted by a classifier trained on the original dataset.

#### 3.1 Generating Domain Descriptions

A key idea of our method is to use captioning and language models trained on large amounts of pretraining data to summarize the potential domains. We assume knowledge of the superclass that

encapsulates all classes, such as “bird” for CUB or “animal” for iWildCam and a prefix to guide the format of the desired caption form, for example “a photo of a bird...”. Once these prompts are generated, we found that we often achieve higher quality edits by adding the class name into the prompts after the fact (e.g. “a photo of a Scott Oriole bird...”).

To generate a concise list of domain descriptions, we first caption each image in the dataset using a pretrained captioning model. This produces a comprehensive set of captions, which may highlight potential domains seen at test time. Note that these captions do not need to accurately describe the task-specific information, such as the species of the bird, as their purpose is to provide a broad overview of the context, such as the environment the bird is in or the actions of the bird. Additional image data of possible domains seen in testing but that don’t contain any task-relevant information can also be used in this step. For example, when performing animal classification, one can add in background images of different locations that may be seen in deployment.

We assume that the amount of training data is not small ( $< 100$  samples). Therefore, we use an LLM to summarize these captions into a list of domains which are agnostic of the class. Due to constraints on context length, we randomly sample 200 unique captions and use the following prompt:

“I have a set of image captions that I want to summarize into objective descriptions that describe the scenes, actions, camera pose, zoom, and other image qualities present. My captions are [CAPTIONS]. I want the output to be a handful of captions that describe a unique setting, of the form [PREFIX]”

We then ask a refinement question to ensure each caption is of only one setting and agnostic of the class with the following prompt:

“Can you modify your response so each caption is agnostic of the type of [SUPER-CLASS]. Please output less than 10 captions which cover the largest breadth of concepts.”

The end result is a list of less than 10 domain descriptions that often include domains seen in the test data (e.g. “A camera trap photo of an animal near a lake.”). These descriptions serve as the foundation for the subsequent text-conditioned image editing stage of our method. We use BLIP [19] captioning model and GPT-4 [26] for summarizing the captions. A complete list of prompts for each dataset is given in Section 5.

### 3.2 Editing Images with Language Guidance

Once the domain descriptions are generated, they are used to condition the edits of original images in the training set. In our experiments, we employ two editing techniques based on Stable Diffusion; however, ALIA can be used with any text-conditioned image editing method. The two techniques deployed for our experiments are as follows:

*Image to Image with Text Guidance (Img2Img) [31, 2, 23]:* This technique first uses an image encoder to translate a provided image into a latent representation. Then, leveraging a diffusion model, this latent representation is progressively modified through a series of transformations, conditioned on the user-provided text. Finally, the modified latent representation is decoded to generate an augmented image that incorporates the desired modifications specified in the prompt.

*Instruct Pix2Pix [3]:* Given an edit instruction (e.g. “put the animals in the forest”) and an image, this method generates a modified image adhering to the specified instruction. This is accomplished by training a conditional diffusion model on a dataset composed of paired images and edit instructions.

Among the many existing image editing techniques [24, 11, 4, 44, 27], we selected the two above for their ease of use and quality of outputs. Notably, the methods highlighted here share a common goal: to generate coherent, visually-grounded augmentations conditioned on the natural language descriptions extracted from the dataset.

### 3.3 Filtering Failed Edits Using Semantic and Visual Features

As depicted in Figure 3, there are three distinct failure cases for the text-conditioned augmentation: (1) total failure, where the edited image is vastly different from the original, (2) identity failure, where



Figure 3: **Filtering.** Semantic filtering eliminates instances of total failure (left), but often misses instances of minimal edits (center) or edits which corrupt class information (right). Meanwhile, our confidence-based filtering mechanism is able to remove these failures by leveraging the prediction confidence of classifier trained on the original dataset applied to the image edits.

the edited image is nearly identical to the original, and (3) class corruption failure, where the edit significantly changed task-specific features. While previous work [9] utilized CLIP-based filtering to remove low-quality images, it only removes instances of total failure in settings where CLIP does not have a good representation of the classes. As such, we also employ a confidence-based filtering technique which uses a classifier trained on the original training set to determine instances of identity and class corruption failures.

*Semantic Filtering.* We use CLIP to predict whether the generated image is related to the task or not. For example, in the CUB dataset, we provide the text prompt “a photo of a bird” as well as the filtering prompts “a photo of an object”, “a photo of a scene”, “a photo of geometric shapes”, “a photo”, and “an image”. All images that are not classified as “a photo of a bird” are removed.

*Confidence-based Filtering.* We take inspiration from previous work in identifying mislabeled examples using model confidence [25]. After training a model  $f$  on the original dataset, we calculate a confidence threshold  $t_y$  for each class  $y$  by averaging the softmax score of the correct label for each image in the training set. Specifically, given edited image  $x'$  with label  $y$  and prediction  $\hat{y}$ , it is filtered out if  $\text{confidence}(f(x'), \hat{y}) \geq t_{\hat{y}}$ . In the case that the image was correctly predicted ( $y = \hat{y}$ ), since the model is already confident in its prediction, the edited image is likely to contain roughly the same information as the original, and thus should be filtered out. In cases where  $y \neq \hat{y}$ , high confidence suggests that the edit has corrupted the class to a point that it more highly resembles another class.

In our experiments, we randomly sample a portion of the filtered augmentations to incorporate back into the training set, leading to a dataset size expansion between 20-100%. Given that the prompts and the image edits are grounded in the original training set, ALIA is able to preserve visual consistency and encompass a broader array of domains.

## 4 Experimental setup

### 4.1 Implementation.

We fine-tune a ResNet50 for the CUB [41] and iWildCam [17] datasets and a ResNet18 for the Planes [22] dataset. We use the PyTorch pretrained models [28] on ImageNet with an Adam optimizer [16] and cosine learning rate scheduler. For each method, we do a hyperparameter sweep across learning rate and weight decay and choose the parameters with the highest validation performance. We train on 10 GeForce RTX 2080 Ti GPUs.

For all the diffusion-based editing methods, we use Stable Diffusion version 1.5 [32] from HuggingFace [40] with the default hyperparameters aside from the edit strength (how much to deviate from the original image) and the text guidance (how closely the generated image should align with the text prompt). For these parameters, we search over 5 different values each for edit strength and text guidance, visualizing the resulting generations for a random sample (10 images) across 4 random seeds (40 generated images in total). We pick the parameters which generate the most diverse images that both retain the task-relevant information and remain faithful to the edit instruction. More details on this selection process including visualizations as well as results on training on data generated with different edit strengths and text guidance are in the Appendix. Results are averaged over 3 random seeds and further details on the hyperparameter search space and final choice of hyperparameters are listed in the Appendix.

## 4.2 Baselines.

The *Baseline* model is trained on the original training dataset, without any data augmentation. We additionally compare adding in the generated data with our method to the original training set (+ALIA), adding in real data from the test distribution (+Real), adding in diffusion generated data from text alone (+Txt2Img), and two traditional augmentation baselines (+CutMix, +RandAug). In order to perform a fair comparison, we keep the number of images added per class to the training set consistent across methods.

For the +ALIA and +Txt2Img results, we generate twice as much data as the original training set for each language prompt to ensure enough data is available after filtering. For +ALIA, this entails generating 2 edits for every image. We then randomly sample from these datasets to match the class distribution of the held-out real data used for the +Real baseline. For the +Txt2Img baseline, we generate images with a hand-crafted prompt for each dataset, and apply our semantic filtering to the generated images to remove low-quality samples. We provide the amount of data added and the prompts used for the +Txt2Img baseline in the subsequent sections.

Lastly, we compare to two data augmentation baselines taken from recent literature: (1) *CutMix* [42], which generates mixed samples by randomly cutting and pasting patches between training images to encourage the model to learn more localized and discriminative features and (2) *RandAugment* [5], an automated data augmentation technique that reduces the search space of the type and magnitude of augmentations to find the best image transformations for a given task. For these baselines, we use the implementations in the PyTorch Torchvision library [21].

## 5 Experiments

We evaluate on 3 specialized tasks: domain generalization via camera trap animal classification (iWildCam), fine-grained bird classification (CUB), and fine-grained aircraft classification with contextual bias (Airbus VS Boeing). Details of each dataset are listed in subsequent sections, with more statistics in the Appendix.

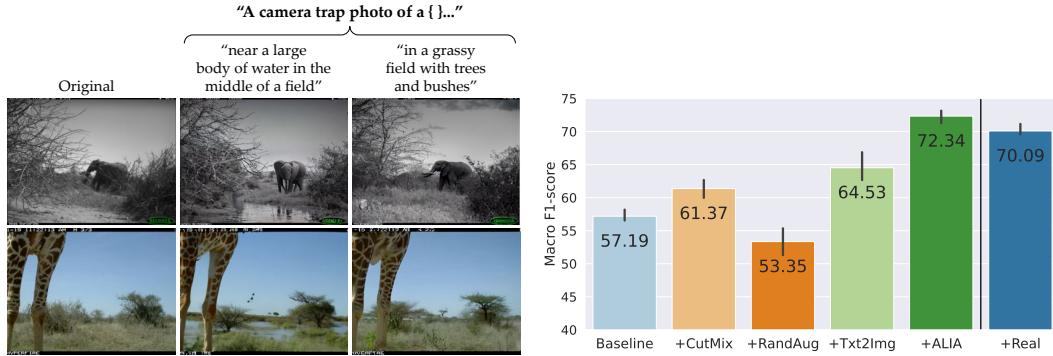


Figure 4: **iWildCam Subset.** Left plot visualizes the original training images and the corresponding data generated from ALIA based on the generated prompts. Right plot shows macro F1-scores for adding in generated data from ALIA and baselines as in Section 4.2. ALIA significantly outperforms all the baselines including adding in real data from the test distribution, while +Txt2Img and +CutMix see smaller improvements and +RandAug results in lower performance than the baseline.

### 5.1 Domain Generalization [iWildCam [17]]

The iWildCam dataset is a large-scale collection of images captured from camera traps placed in various locations around the world. We subsample the dataset to create a 7-way classification task (background, cattle, elephant, impala, zebra, giraffe, dik-dik), with 2 test locations that are not in the training or validation set. The training set has 6,000 images with some classes having as few as 50 images per example, and our +Real data baseline contains roughly 2,000 images from locations not seen during training. We generate our domain descriptions from the background images from the test domain, but do not use these images in training. Since many images in the original training set

contain watermarks or timestamps while the text-to-image generated data usually does not, we crop the top and bottom of each image in prepossessing.

The prompts generated by ALIA are “*a camera trap photo of a { }...*”:

- (1) “*in a grassy field with trees and bushes.*”, (2) “*in a forest in the dark.*”, (3) “*near a large body of water in the middle of a field.*”, (4) “*walking on a dirt trail with twigs and branches.*”.

We use “*a camera trap photo of a { } in the wild.*” as the prompt for the *+Txt2Img* baseline.

As shown in Figure 4, ALIA not only outperforms all baselines, with a 15% performance improvement in f1-score over training on the original data, but even exceeds the performance of adding in the same amount of real data. Note the improved performance over real can be somewhat attributed to that fact that our prompts are taken from empty images from the test domain rather than the training set.

## 5.2 Fine-grained Classification [CUB [41]]

CUB is a fine-grained bird classification dataset comprised of photos taken from Flickr. We remove 5 of the 30 images from each of the 200 classes from the training set for the *+Real* comparison.

The prompts generated by ALIA are “*a photo of a { } bird...*”:

- (1) “*interacting with flowers.*”, (2) “*standing by the waters edge.*”, (3) “*perched on a fence.*”, (4) “*standing on a rock.*”, (5) “*perched on a branch.*”, (6) “*flying near a tree, sky as the backdrop.*”, (7) “*perched on a birdfeeder.*”

We use “*an iNaturalist photo of a { } bird in nature.*” as the prompt for the *+Txt2Img* baseline.

As shown in Figure 5, ALIA outperforms all the baselines aside from adding in real data, while adding text-to-image generated data results in similar performance with the baseline. Significantly, these results show that ALIA provides performance improvements greater than existing data augmentation techniques even in conditions devoid of domain shifts, thereby broadening its utility.

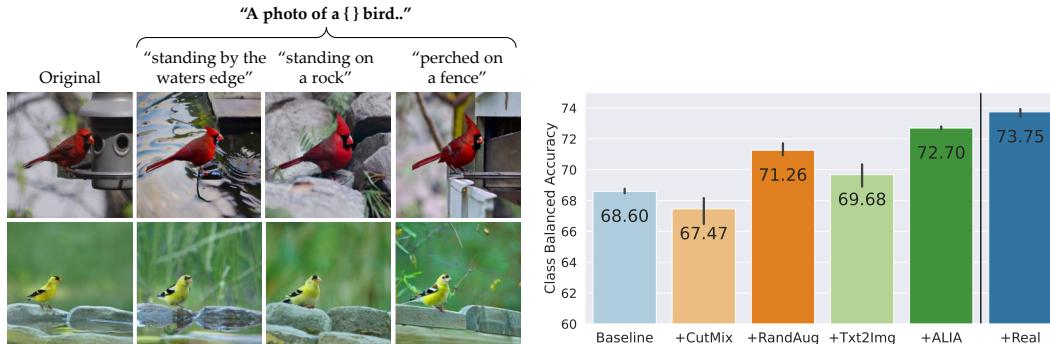
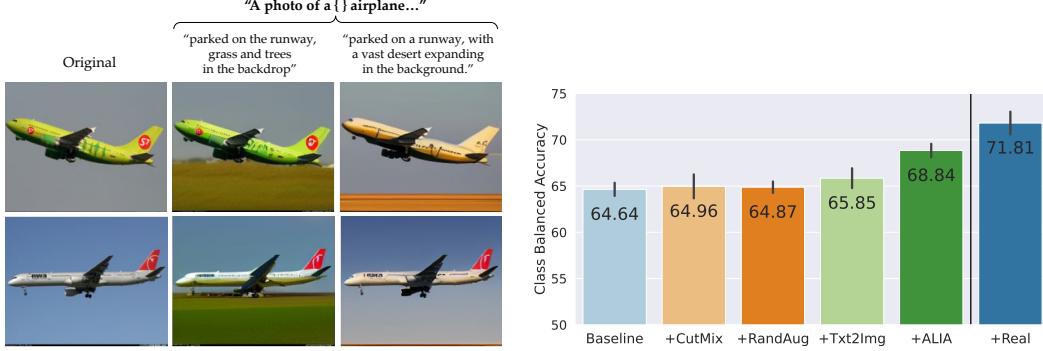


Figure 5: **Bird Classification (CUB).** Left plot visualizes the original training images and the corresponding data generated from ALIA based on the generated prompts. Right plot shows class-balanced accuracy for ALIA and baselines. ALIA outperforms all the baselines aside from adding in real data, while adding text-to-image generated data results in similar performance with the baseline.

## 5.3 Contextual Bias [Airbus VS Boeing [22]]

In order to test our method on a setting of contextual bias for fine-grained classification, we create a custom split of the FGVC-Aircraft Benchmark [22] containing 2 visually similar classes: Boeing-767 and Airbus-322. We then manually label each image as ‘sky’, ‘grass’, or ‘road’ based on the background and filter out ambiguous examples. To construct a bias split, we train on 400 samples where Airbus airplanes are only seen in road backgrounds and Boeing airplanes are only seen on grassy backgrounds, with both appearing in sky backgrounds. We evaluate on a testset where each class appears on all backgrounds. Similarly, our *+Real* baseline adds 400 samples of a similar background distribution to the test set. Exact split breakdowns are in the Appendix.

Interestingly, we found that the edits generated by the Img2Img editing method were almost exclusively total failures (visualizations in the Appendix). Thus, we use InstructPix2Pix, prompting the



**Figure 6: Airbus VS Boeing Classification.** Left plot visualizes the original training images and the corresponding data generated from ALIA based on the generated prompts. Right plot shows class-balanced accuracy for ALIA and baselines. ALIA outperforms all the baselines aside from adding in real data.



**Figure 7: Examples of filtered edits of AirBus VS Boeing with confidence-based filtering.** Our confidence-based filtering effectively prunes overrepresented class-domain pairings (e.g., Airbus on road, Boeing on grass) and boosts the prevalence of underrepresented ones (e.g., Airbus on grass, Boeing on road), thereby enhancing the diversity and balance in the augmented data.

LLM with the prefixes “a photo of an airplane” and editing the final prompts to fit the instructional format of the editing method.

The prompts generated by ALIA are “*a photo of a {} airplane...*”:

- (1) “*on the airport tarmac, surrounded by buildings and other infrastructure.*”, (2) “*parked on the runway, grass and trees in the backdrop.*”, (3) “*parked on a runway, with a vast desert expanding in the background.*”, (4) “*with red and white colors, landing gear down, against a backdrop of a bustling cityscape.*”, (5) “*in mid-flight, landing gear deployed against a clear sky.*”

We use “*a photo of a {} airplane.*” as the prompt for the +Txt2Img baseline.

The prompts generated by ALIA are not only capable of identifying the backgrounds used to create this biased split, but as illustrated in Figure 7, our confidence-based filtering technique can effectively remove many instances of the overrepresented class-domain pairings (Airbus on road and Boeing on grass) from the augmented dataset. These instances are well-represented in the training data, and consequently, the classifier can predict them with high confidence. This results in a higher proportion of the edited data consisting of the underrepresented classes (Airbus on grass and Boeing on road). As a result of these two factors, ALIA outperforms all other augmentation baselines on a test set where each class is seen in the sky, on the grass, and on the road (Figure 6). We believe that these results point to a promising future direction in how to modify ALIA to identify and amplify specific class-domain pairings which are underrepresented in the training set.

Dataset	Metric	User Prompt	ALIA Prompts	ALIA Prompts + Filtering
iWildCam	Macro F1-score	$68.87 \pm 1.84$	$70.65 \pm 1.50$	$72.34 \pm 1.00$
CUB	Balanced Accuracy	$71.02 \pm 0.47$	$71.25 \pm 0.86$	$72.70 \pm 0.10$
Planes	Balanced Accuracy	$62.453 \pm 1.03$	$67.03 \pm 0.65$	$68.84 \pm 0.89$

**Table 1: Effect of ALIA Generated Prompts and Filtering.**

## 5.4 Ablations

**Effect of Prompt Quality and Filtering.** We ablate the quality of ALIA generated prompts by comparing the edit prompts generated by ALIA against those provided by users. For the user-supplied prompts, we draw on the engineered prompts that were used for the  $+Txt2Img$  baseline. As shown in Table 1, descriptions generated by ALIA outperform user-provided prompts, especially in the contextual bias setting, indicating that our generated prompts are able to accurately describe key domain-specific features and variations, resulting in more effective data augmentation for mitigating biases and improving model generalization. Moreover, Table 1 also demonstrates how our semantic and confidence-based filtering further improves accuracy.

**Choice of Image Editing Method.** As mentioned in Section 5.3, the suitability of image editing methods can vary based on the specific dataset. As such, we explore the impact of the editing method on performance by modifying the iWildCam subset from Section 5.1 with InstructPix2Pix. We employ the same prompts as in the Img2Img case, reformatted to the instructional syntax required by InstructPix2Pix (e.g. “a camera trap photo of a {} in a grassy field with trees and bushes” becomes “put the {} in a grassy field with trees and bushes”). Figure 8 presents examples of edits generated using InstructPix2Pix. Although InstructPix2Pix can produce edits with more significant color variations than Img2Img, it frequently applies an unrealistic airbrush-like effect or erases most of the information present in the original image. This outcome can be attributed to the fact that InstructPix2Pix is fine-tuned on a dataset of paired images and edits, which primarily consist of stylistic changes to artistic images, making the iWildCam images fall outside its domain. When comparing the performance of InstructPix2Pix, we find that it achieves a macro F1-score of  $67.11 \pm 1.96$ , compared to  $73.34 \pm 1.0$  obtained using Img2Img edits. These results indicate that the choice of image editing method has significant impact on the final performance of ALIA, and thus it is crucial to profile the quality of image edits from various techniques on one’s specific dataset.



Figure 8: **InstructPix2Pix edits on iWildCam.** Despite offering more color variation, the method often imparts an artificial airbrush-like effect or removes substantial original image details.

## 6 Limitations

While ALIA is able to generate impressive augmentations, there are limitations. As ALIA depends on large pretrained vision and language models to translate between task-specific image data and task-agnostic natural language descriptions of domains, performance of the method is bottlenecked by the quality of the captioning model, LLM, and image editing method (see Section 5.4). Furthermore, as we assume that the task-specific information in the training data cannot be easily generated via text alone, attempts to change aspects like the pose of a bird in CUB is likely to result in a failed edit. Finally, determining the optimal quantity of augmented data to reincorporate into the training set remains an unresolved question, and is an area we look forward to addressing in future research.

## 7 Conclusion

We present ALIA, a novel approach to data augmentation that leverages the high-level domain knowledge of large language models and text-conditioned image editing methods. By grounding both the domain descriptions and the augmented training data in the provided training set, our method has demonstrated impressive capabilities in several challenging settings, including domain adaptation, bias mitigation, and even scenarios without domain shift. As the capabilities of captioning, LLMs, and image editing methods grow, we expect the efficacy and scope of our approach to increase.

**Acknowledgements.** We thank Suzie Petryk for her invaluable feedback on the manuscript. This work was supported in part by the NSF CISE Expeditions Award (CCF-1730628), DARPA’s SemaFor, PTG and/or LwLL programs, and BAIR’s industrial alliance programs.

## References

- [1] Shekoofeh Azizi, Simon Kornblith, Chitwan Saharia, Mohammad Norouzi, and David J. Fleet. Synthetic data from diffusion models improves imagenet classification, 2023.
- [2] Andreas Blattmann, Robin Rombach, Kaan Oktay, and Björn Ommer. Retrieval-augmented diffusion models, 2022.
- [3] Tim Brooks, Aleksander Holynski, and Alexei A. Efros. Instructpix2pix: Learning to follow image editing instructions. November 2022.
- [4] Guillaume Couairon, Jakob Verbeek, Holger Schwenk, and Matthieu Cord. Diffedit: Diffusion-based semantic image editing with mask guidance. *arXiv preprint arXiv:2210.11427*, 2022.
- [5] Ekin Dogus Cubuk, Barret Zoph, Jon Shlens, and Quoc Le. Randaugment: Practical automated data augmentation with a reduced search space. In H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 18613–18624. Curran Associates, Inc., 2020.
- [6] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- [7] Lisa Dunlap, Clara Mohri, Han Zhang, Devin Guillory, Trevor Darrell, Joseph E. Gonzalez, Anna Rohrbach, and Aditi Raghunathan. Using language to extend to unseen domains. *International Conference on Learning Representations (ICLR)*, 2023.
- [8] Maayan Frid-Adar, Idit Diamant, Eyal Klang, Michal Amitai, Jacob Goldberger, and Hayit Greenspan. Gan-based synthetic medical image augmentation for increased cnn performance in liver lesion classification. *Neurocomputing*, 2018.
- [9] Ruifei He, Shuyang Sun, Xin Yu, Chuhui Xue, Wenqing Zhang, Philip Torr, Song Bai, and Xiaojuan Qi. Is synthetic data from generative models ready for image recognition? *arXiv preprint arXiv:2210.07574*, 2022.
- [10] Dan Hendrycks, Norman Mu, Ekin D. Cubuk, Barret Zoph, Justin Gilmer, and Balaji Lakshminarayanan. AugMix: A simple data processing method to improve robustness and uncertainty. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2020.
- [11] Amir Hertz, Ron Mokady, Jay Tenenbaum, Kfir Aberman, Yael Pritch, and Daniel Cohen-Or. Prompt-to-prompt image editing with cross attention control. *arXiv preprint arXiv:2208.01626*, 2022.
- [12] Sheng-Wei Huang, Che-Tsung Lin, Shu-Ping Chen, Yen-Yi Wu an Po-Hao Hsu, and Shang-Hong Lai. Auggan: Cross domain adaptation with gan-based data augmentation. *European Conference on Computer Vision*, 2018.
- [13] Saachi Jain, Hannah Lawrence, Ankur Moitra, and Aleksander Madry. Distilling model failures as directions in latent space. In *ArXiv preprint arXiv:2206.14754*, 2022.
- [14] Jang-Hyun Kim, Wonho Choo, and Hyun Oh Song. Puzzle mix: Exploiting saliency and local statistics for optimal mixup. In *International Conference on Machine Learning (ICML)*, 2020.
- [15] Jang-Hyun Kim, Wonho Choo, and Hyun Oh Song. Puzzle mix: Exploiting saliency and local statistics for optimal mixup. In *International Conference on Machine Learning*, pages 5275–5285. PMLR, 2020.
- [16] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [17] Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanas Phillips, Irena Gao, et al. Wilds: A benchmark of in-the-wild distribution shifts. In *International Conference on Machine Learning*, pages 5637–5664. PMLR, 2021.
- [18] Daiqing Li, Huan Ling, Seung Wook Kim, Karsten Kreis, Adela Barriuso, Sanja Fidler, and Antonio Torralba. Bigdatasetgan: Synthesizing imagenet with pixel-wise annotations, 2022.
- [19] Junnan Li, Dongxu Li, Caiming Xiong, and Steven Hoi. Blip: Bootstrapping language-image pre-training for unified vision-language understanding and generation, 2022.
- [20] Zicheng Liu, Siyuan Li, Di Wu, Zihan Liu, Zhiyuan Chen, Lirong Wu, and Stan Z Li. Automix: Unveiling the power of mixup for stronger classifiers. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XIV*, pages 441–458. Springer, 2022.
- [21] TorchVision maintainers and contributors. Torchvision: Pytorch’s computer vision library. <https://github.com/pytorch/vision>, 2016.
- [22] S. Maji, J. Kannala, E. Rahtu, M. Blaschko, and A. Vedaldi. Fine-grained visual classification of aircraft. Technical report, 2013.
- [23] Chenlin Meng, Yutong He, Yang Song, Jiaming Song, Jiajun Wu, Jun-Yan Zhu, and Stefano Ermon. SDEdit: Guided image synthesis and editing with stochastic differential equations. In *International Conference on Learning Representations*, 2022.
- [24] Alex Nichol, Prafulla Dhariwal, Aditya Ramesh, Pranav Shyam, Pamela Mishkin, Bob McGrew, Ilya Sutskever, and Mark Chen. Glide: Towards photorealistic image generation and editing with text-guided diffusion models. In *ICML*, 2022.

- [25] Curtis G. Northcutt, Anish Athalye, and Jonas Mueller. Pervasive label errors in test sets destabilize machine learning benchmarks. In *Proceedings of the 35th Conference on Neural Information Processing Systems Track on Datasets and Benchmarks*, December 2021.
- [26] OpenAI. Gpt-4 technical report, 2023.
- [27] Dong Huk Park, Grace Luo, Clayton Toste, Samaneh Azadi, Xihui Liu, Maka Karalashvili, Anna Rohrbach, and Trevor Darrell. Shape-guided diffusion with inside-outside attention. *arXiv*, 2022.
- [28] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, et al. Pytorch: An imperative style, high-performance deep learning library. *Advances in neural information processing systems*, 32, 2019.
- [29] William Peebles, Jun-Yan Zhu, Richard Zhang, Antonio Torralba, Alexei Efros, and Eli Shechtman. Gan-supervised dense visual alignment. In *CVPR*, 2022.
- [30] Suzanne Petryk, Lisa Dunlap, Keyan Nasseri, Joseph Gonzalez, Trevor Darrell, and Anna Rohrbach. On guiding visual attention with language specification. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- [31] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models, 2021.
- [32] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *CVPR*, 2022.
- [33] Swami Sankaranarayanan, Yogesh Balaji, Carlos Domingo Castillo, and Rama Chellappa. Generate to adapt: Aligning domains using generative adversarial networks. *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
- [34] Viktoriia Sharmanska, Lisa Anne Hendricks, Trevor Darrell, and Novi Quadrianto. Contrastive examples for addressing the tyranny of the majority. *CoRR*, abs/2004.06524, 2020.
- [35] Jordan Shipard, Arnold Wiliem, Kien Nguyen Thanh, Wei Xiang, and Clinton Fookes. Boosting zero-shot classification with synthetic data diversity via stable diffusion. *arXiv preprint arXiv:2302.03298*, 2023.
- [36] Connor Shorten and Taghi M Khoshgoftaar. A survey on image data augmentation for deep learning. *Journal of big data*, 6(1):1–48, 2019.
- [37] Rohan Taori, Achal Dave, Vaishaal Shankar, Nicholas Carlini, Benjamin Recht, and Ludwig Schmidt. Measuring robustness to natural distribution shifts in image classification. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [38] Brandon Trabucco, Kyle Doherty, Max Gurinas, and Ruslan Salakhutdinov. Effective data augmentation with diffusion models, 2023.
- [39] Nontawat Tritrong, Pitchaporn Rewatbowornwong, and Supasorn Suwanjanakorn. Repurposing gans for one-shot semantic part segmentation. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.
- [40] Patrick von Platen, Suraj Patil, Anton Lozhkov, Pedro Cuenca, Nathan Lambert, Kashif Rasul, Mishig Davaadorj, and Thomas Wolf. Diffusers: State-of-the-art diffusion models. <https://github.com/huggingface/diffusers>, 2022.
- [41] Catherine Wah, Steve Branson, Peter Welinder, Pietro Perona, and Serge Belongie. The caltech-ucsd birds-200-2011 dataset, 2011.
- [42] Sangdoo Yun, Dongyoon Han, Seong Joon Oh, Sanghyuk Chun, Junsuk Choe, and Youngjoon Yoo. Cutmix: Regularization strategy to train strong classifiers with localizable features. In *International Conference on Computer Vision (ICCV)*, 2019.
- [43] Hongyi Zhang, Moustapha Cissé, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. *International Conference on Learning Representations*, 2018.
- [44] Lvmin Zhang and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models, 2023.
- [45] Yuxuan Zhang, Huan Ling, Jun Gao, Kangxue Yin, Jean-Francois Lafleche, Adela Barriuso, Antonio Torralba, and Sanja Fidler. Datasetgan: Efficient labeled data factory with minimal human effort. In *CVPR*, 2021.

## Supplimentary Material

### 8 Dataset Statistics

We provide the number of images for each dataset split used in Section 5 in Table 2, where Extra denotes the amount of data added for the real data baseline. This extra set was constructed for each dataset manually. Furthermore, we provide the number of images for each background class (sky, grass, road) used to create a biased training set used in Section 5.3 in Table 3.

Dataset	Training	Extra	Validation	Testing	% added
iWildCam	6052	2224	2826	8483	37%
CUB	4994	1000	5794	5794	20%
Airbus VS Boeing	409	357	358	707	87%

Table 2: Dataset Split Sizes. Note that CUB’s validation set and test set are the same.

	Airbus			Boeing		
	sky	grass	road	sky	grass	road
Train	98	0	70	129	112	0
Extra	90	21	21	136	44	45
Val	90	21	21	137	45	44
Test	175	51	51	222	104	104

Table 3: Dataset Statistics for Planes

### 9 Hyperparameters

To select the hyperparameters, we train the baseline with learning rates of [0.00001, 0.0001, 0.001, 0.01, 0.1] and weight decay [1e-5, 1e-4, 1e-3, 1e-2], selecting the configuration that results in the highest validation accuracy. These parameters, shown in Table 4, are then used across all methods.

Dataset	Training			Img2Img		InstructPix2Pix	
	learning rate	weight decay	epochs	edit strength	text guidance	edit strength	text guidance
iWildCam	0.0001	1e-4	100	0.4	5.0	1.3	7.5
CUB	0.001	1e-4	200	0.6	7.5	-	-
Airbus VS Boeing	0.001	1e-4	200	0.3	5.0	1.3	7.5

Table 4: Hyperparameters

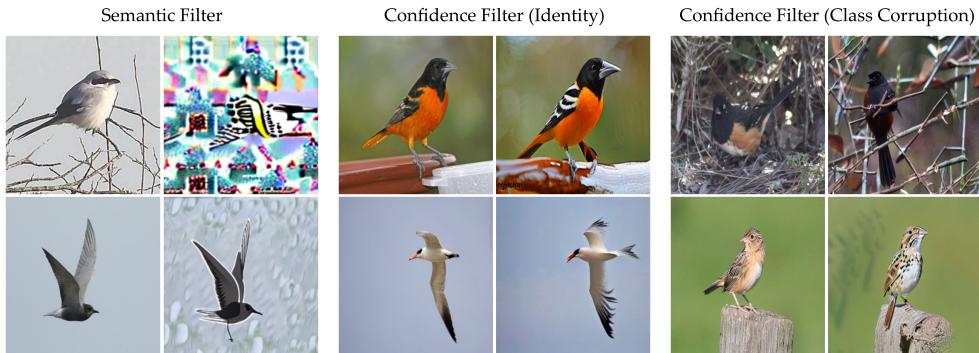


Figure 9: Random Sampled Edits filtered from Cub.

## 10 Additional Visualizations

### 10.1 Filtered Images

We present qualitative evidence confirming that our semantic and confidence filters effectively remove failed edits. Figure 9 showcases edits filtered from Cub with the semantic filter (left), the confidence filter with high confidence and correct prediction (center), and the confidence filter with high confidence but incorrect prediction (right). As detailed in Section 3.3, the semantic filter removes total failures or irregular looking images, while the confidence-based filter eliminates edits that are excessively similar to the original image or that corrupt task-relevant information. Edits were created using Img2Img with the prompt “a photo of a bird perched on a branch.”

### 10.2 Different Strengths and Text Guidance

The performance of ALIA in Section 5 is significantly influenced by the hyperparameters related to the image editing techniques, namely, the edit strength and text guidance.

The edit strength hyperparameter in the stable diffusion model controls the extent of modifications to the original image, with higher values yielding more noticeable and significant changes. We edited iWildCam images using the prompt “a cameratrap photo of a near a large body of water in the middle of a field,” employing Img2Img and InstructPix2Pix with edit strength ranging from  $0.1 \rightarrow 0.9$  for Img2Img and  $1.1 \rightarrow 1.9$  for InstructPix2Pix. Both Img2Img and InstructPix2Pix had text guidance set to 7.5. As illustrated in Figure 10, increasing the edit strength leads to greater departure from the original image, albeit at the risk of class information corruption.

The text guidance hyperparameter in the stable diffusion model determines how closely the generated output aligns with the textual description. Higher values produce outputs that more explicitly conform to the provided text. We edited iWildCam images using the prompt “a cameratrap photo of a near a large body of water in the middle of a field,” using Img2Img and InstructPix2Pix with text guidance values of 5.0, 7.5, and 9.0. The edit strength was set to 0.4 for Img2Img and 1.3 for InstructPix2Pix. As Figure 11 demonstrates, higher text guidance results in images with more evident water, as specified in the prompt. However, similar to the edit strength hyperparameter, higher text guidance may also lead to greater class information corruption.

In practice, we found that selecting strength and text guidance values that yield more pronounced edits ultimately resulted in higher performance. This is because such edits significantly alter the context, while the filtering step mitigates the effects of detrimental edits. For example, running the same experiment as described in Section 5.1 with strength = 0.3 and text guidance = 5.0 resulted in a macro F1-score of 64.25. This marks an approximate 8 point drop compared to strength = 0.6 and text guidance = 7.5 (72.34).

### 10.3 Generated Images

Figures 12, 13, and 14 show examples from each prompt used to create the ALIA augmented dataset (Section 5), before the filtering step.

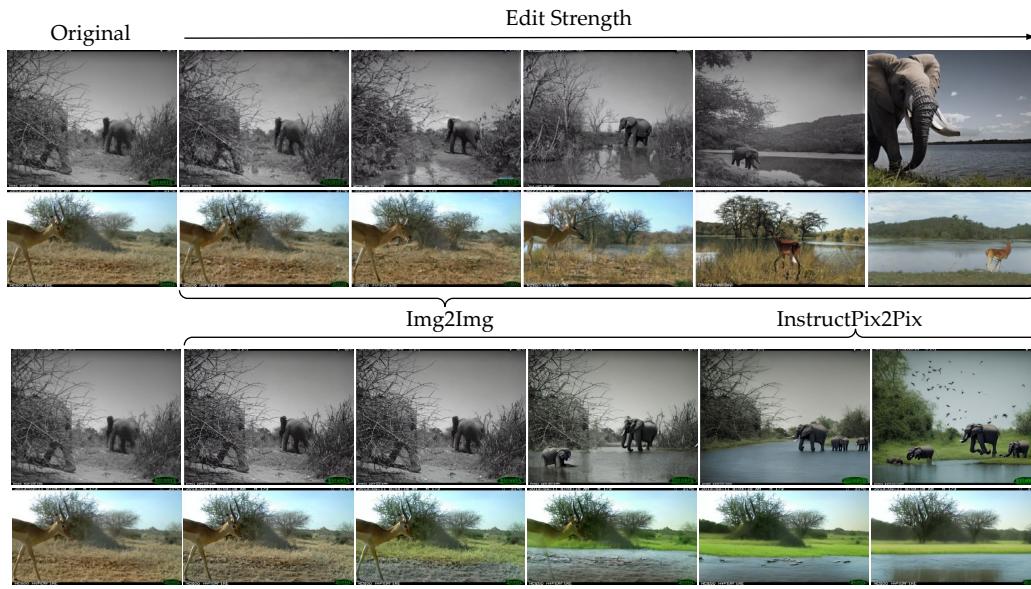


Figure 10: Different Edit Strengths.

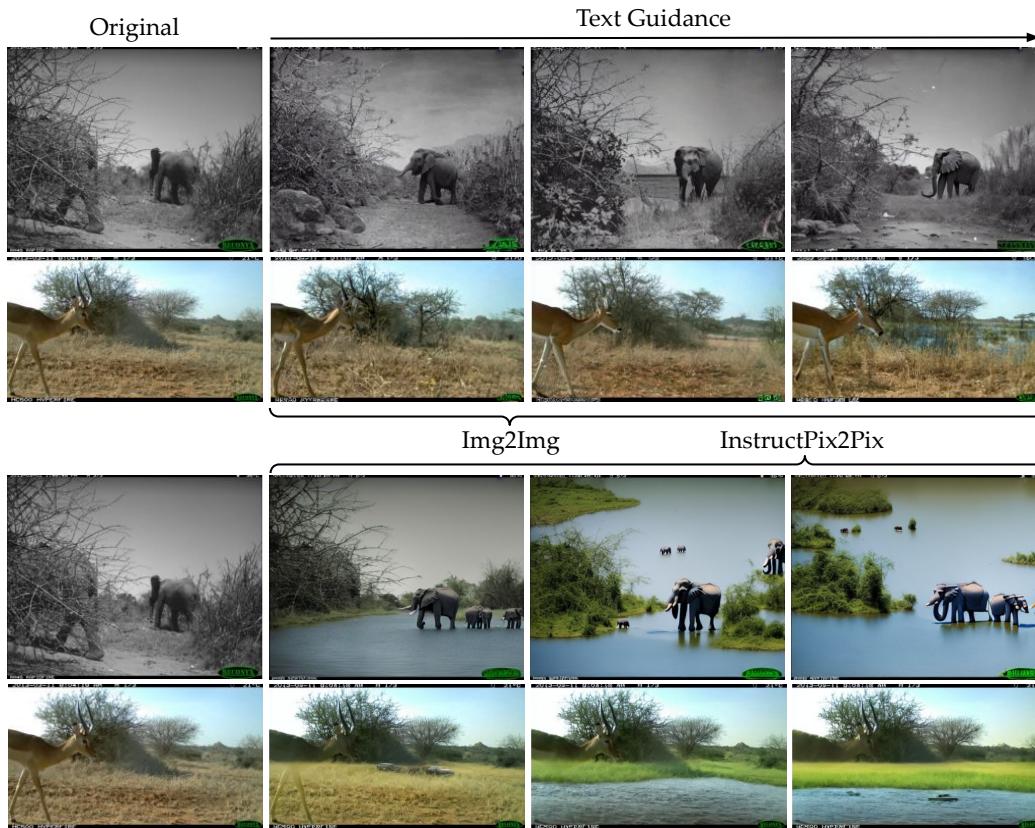
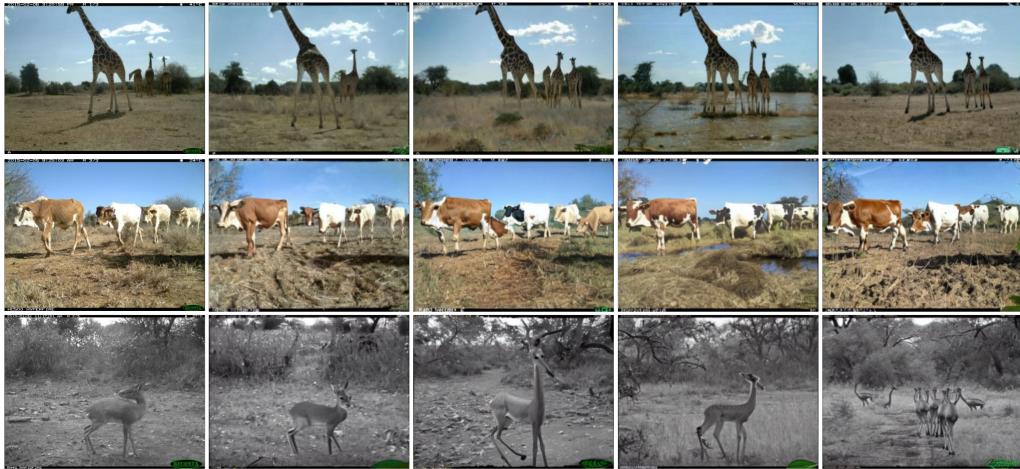
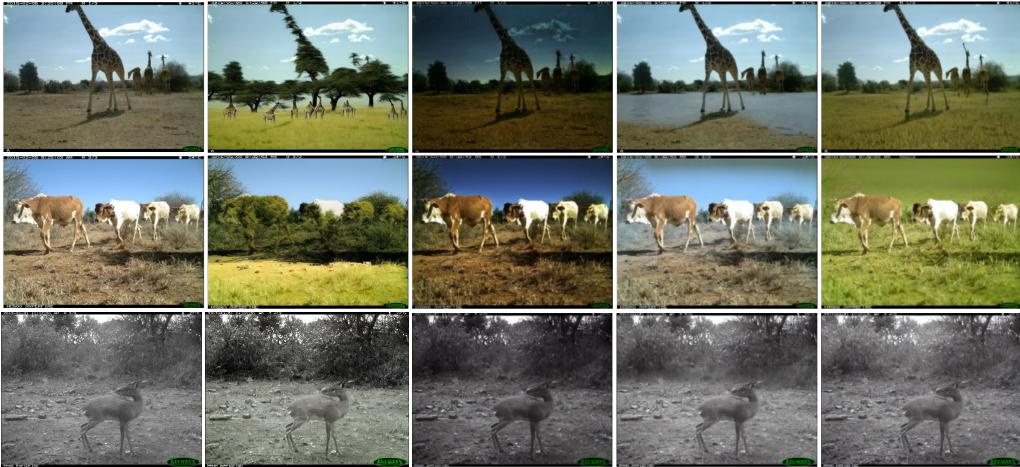


Figure 11: Different Text Guidance.

Img2Img



InstructPix2Pix



Txt2Img

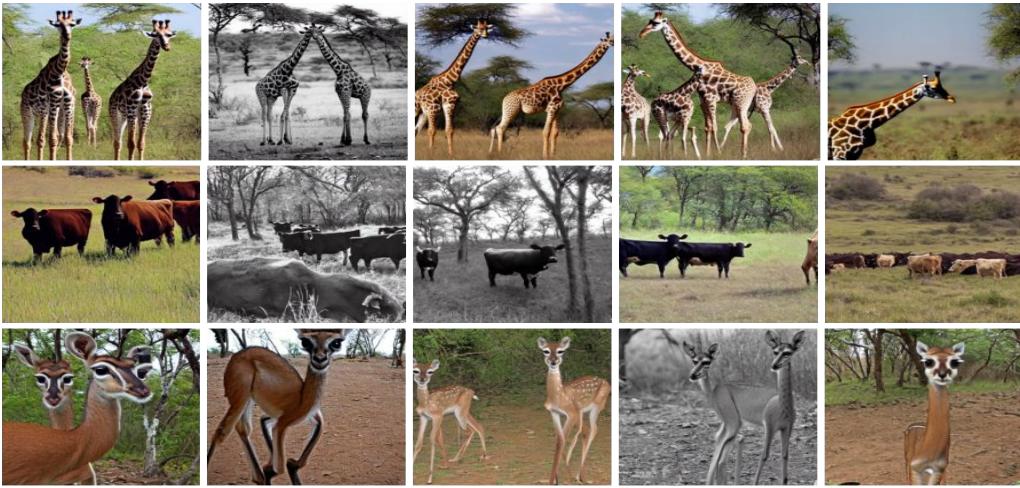


Figure 12: Randomly sampled edits of iWildCam, unfiltered. For Img2Img, the first column corresponds to the original image, and the further columns correspond to the following prompts (from left to right): “*a camera trap photo of a { }...*”: (1) “*in a grassy field with trees and bushes.*”, (2) “*in a forest in the dark.*”, (3) “*near a large body of water in the middle of a field.*”, (4) “*walking on a dirt trail with twigs and branches.*”. Txt2Img samples are generated with the prompt “*a camera trap photo of a { } in the wild.*”.



Figure 13: Randomly sampled edits of Cub, unfiltered. For Img2Img, the first column corresponds to the original image, and the further columns correspond to the following prompts (from left to right): “*a photo of a {} bird...*”: (1) “*interacting with flowers.*”, (2) “*standing by the waters edge.*”, (3) “*perched on a fence.*”, (4) “*standing on a rock.*”, (5) “*perched on a branch.*”, (6) “*flying near a tree, sky as the backdrop.*”, (7) “*perched on a birdfeeder.*”. Txt2Img samples are generated with the prompt “*an iNaturalist photo of a {} bird in the wild.*”.

InstructPix2Pix



Img2Img



Txt2Img



Figure 14: Randomly sampled edits of Airbus (top) and Boeing (bottom), unfiltered. For InstructPix2Pix and Img2Img, the first column corresponds to the original image, and the further columns correspond to the following prompts (from left to right) “*a photo of a { } airplane..*”: (1) “*parked on the runway, grass and trees in the backdrop.*”, (2) “*on the airport tarmac, surrounded by buildings and other infrastructure.*”, (3) “*parked on a runway, with a vast desert expanding in the background.*”, (4) “*with red and white colors, landing gear down, against a backdrop of a bustling cityscape.*”, (5) “*in mid-flight, landing gear deployed against a clear sky.*”. Txt2Img samples are generated with the prompt “*a photo of a { } airplane.*”.